

1. Configuration screening

11 configs × 2 outcomes;
baseline: direct answering,
T=0, first-person framing

2. Sampling & temperature

Shortlisted configs ×
{direct, verbalized}
× $T \in \{0, 0.4, 0.8\}$

3. Model comparison

Leading setups re-run across:
GPT-5.4-mini, Gemini 2.5
Flash Lite, Llama 3.3 70B

4. Presentation stress tests

Address mode, explicit
reasoning instruction,
ideology background paragraph,
natural rewrite

Calibrated protocol

One inference setup per
outcome, fixed before
the main experiment

